

we shall see later. This type of clipping is thus accomplished on the fly; if the bounds check can be done quickly (e.g., by a tight inner loop running completely in microcode or in an instruction cache), this approach may actually be faster than first clipping the primitive and then scan converting the resulting, clipped portions. It also generalizes to arbitrary clip regions.

A third technique is to generate the entire collection of primitives into a temporary canvas and then to copyPixel only the contents of the clip rectangle to the destination canvas. This approach is wasteful of both space and time, but it is easy to implement and is often used for text.

Raster displays invoke clipping and scan-conversion algorithms each time an image is created or modified. Hence, these algorithms not only must create visually satisfactory images but also must execute as rapidly as possible. As discussed in detail in later sections, scan-conversion algorithms use *incremental methods* to minimize the number of calculations (especially multiplies and divides) performed during each iteration; further, these calculations employ integer rather than floating-point arithmetic. Speed can be increased even further by using multiple parallel processors to scan convert simultaneously entire output primitives or pieces of them.

3.2 SCAN CONVERTING LINES

A scan-conversion algorithm for lines computes the coordinates of the pixels that lie on or near an ideal, infinitely thin straight line imposed on a 2D raster grid. In principle, we would like the sequence of pixels to lie as close to the ideal line as possible and to be as straight as possible. Consider a 1-pixel-thick approximation to an ideal line; what properties should it have? For lines with slopes between -1 and 1 inclusive, exactly 1 pixel should be illuminated in each column; for lines with slopes outside this range, exactly 1 pixel should be illuminated in each row. All lines should be drawn with constant brightness, independent of length and orientation, and as rapidly as possible. There should also be provisions for drawing lines that are more than 1 pixel wide, centered on the ideal line, that are affected by line-style and pen-style attributes, and that create other effects needed for high quality illustrations. For example, the shape of the endpoint regions should be under programmer control to allow beveled, rounded, and mitered corners. We would even like to be able to minimize the jaggies due to the discrete approximation of the ideal line, by using antialiasing techniques that exploit the ability to set the intensity of individual pixels on n -bits-per-pixel displays.

For now, we consider only 1-pixel-thick lines that have exactly 1 bilevel pixel in each column (or row for lines with slope $> \pm 1$). Later in the chapter, we will consider thick primitives and deal with styles.

To visualize the geometry, we recall that SRGP represents a pixel as a circular dot centered at that pixel's (x, y) location on the integer grid. This representation is a convenient approximation to the more or less circular cross-section of the CRT's electron beam, but the exact spacing between the beam spots on an actual display

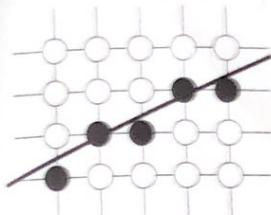


Figure 3.4
A scan-converted line
showing intensified pixels
as black circles.

can vary greatly among systems. In some systems, adjacent spots overlap; in others, there may be space between adjacent vertical pixels; in most systems, the spacing is tighter in the horizontal than in the vertical direction. Another variation in coordinate-system representation arises in systems, such as the Macintosh, that treat pixels as being centered in the rectangular box between adjacent grid lines instead of on the grid lines themselves. In this scheme, rectangles are defined to be all pixels interior to the mathematical rectangle defined by two corner points. This definition allows zero-width (null) canvases: The rectangle from (x, y) to (x, y) contains no pixels, unlike the SRGB canvas, which has a single pixel at that point. For now, we continue to represent pixels as disjoint circles centered on a uniform grid, although we shall make some minor changes when we discuss antialiasing.

Figure 3.4 shows a highly magnified view of a 1-pixel-thick line and of the ideal line that it approximates. The intensified pixels are shown as filled circles, and the nonintensified pixels are shown as unfilled circles. On an actual screen, the diameter of the roughly circular pixel is larger than the interpixel spacing, so our symbolic representation exaggerates the discreteness of the pixels.

Since SRGB primitives are defined on an integer grid, the endpoints of a line have integer coordinates. In fact, if we first clip the line to the clip rectangle, a line intersecting a clip edge may actually have an endpoint with a noninteger coordinate value. The same is true when we use a floating-point raster graphics package. (We will discuss these noninteger intersections in Section 3.2.3.) Assume that our line has slope $|m| \leq 1$; lines at other slopes can be handled by suitable changes in the development that follows. Also, the most common lines—those that are horizontal, are vertical, or have a slope of ± 1 —can be handled as trivial special cases because these lines pass through only pixel centers (see Exercise 3.1).

3.2.1 The Basic Incremental Algorithm

The simplest strategy for scan converting lines is to compute the slope m as $\Delta y / \Delta x$, to increment x by 1 starting with the leftmost point, to calculate $y_i = mx_i + B$ for each x_i , and to intensify the pixel at $(x_i, \text{Round}(y_i))$, where $\text{Round}(y_i) = \text{Floor}(0.5 + y_i)$. This computation selects the closest pixel—that is, the pixel whose distance to the true line is smallest. This brute-force strategy is inefficient, however, because each iteration requires a floating-point (or binary fraction) multiply, addition, and invocation of Floor. We can eliminate the multiplication by noting that

$$y_{i+1} = mx_{i+1} + B = m(x_i + \Delta x) + B = y_i + m\Delta x,$$

and if $\Delta x = 1$, then $y_{i+1} = y_i + m$.

Thus, a unit change in x changes y by m , which is the slope of the line. For all points (x_i, y_i) on the line (*not* the points on our rasterization of the line), we know that if $x_{i+1} = x_i + 1$, then $y_{i+1} = y_i + m$; that is, the values of x and y are defined in terms of their previous values (see Fig. 3.5). This is what defines an incremental algorithm: At each step, we make incremental calculations based on the preceding step.

We initialize the incremental calculation with (x_0, y_0) , the integer coordinates of an endpoint. Note that this incremental technique avoids the need to deal with

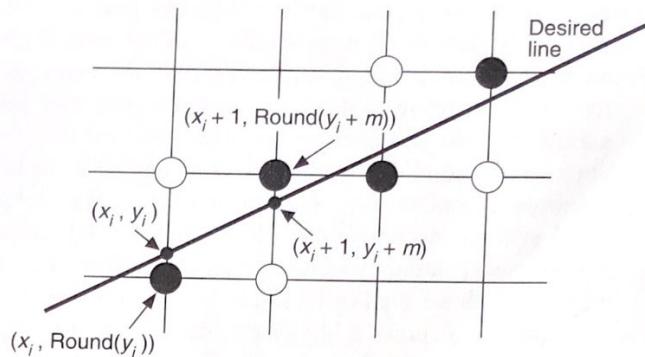


Figure 3.5 Incremental calculation of (x_i, y_i) .

the y intercept, B , explicitly. If $|m| > 1$, a step in x creates a step in y that is greater than 1. Thus, we must reverse the roles of x and y by assigning a unit step to y and incrementing x by $\Delta x = \Delta y/m = 1/m$. Line, the function in Prog. 3.1, implements the incremental technique. The start point must be the left endpoint. Also, it is limited to the case $-1 \leq m \leq 1$, but other slopes may be accommodated by symmetry. Checking for the special cases of horizontal, vertical, or diagonal lines is omitted.

WritePixel, used by Line, is a low-level function provided by the device level software; it places a value into a canvas for a pixel whose coordinates are given as the first two arguments.¹ We assume here that we scan convert only in replace mode; for SRGPs other write modes, we must use a low-level ReadPixel function to read the pixel at the destination location, logically combine that pixel with the source pixel, and then write the result into the destination pixel with WritePixel.

This algorithm is often referred to as a **digital differential analyzer (DDA)** algorithm. The DDA is a mechanical device that solves differential equations by numerical methods: It traces out successive (x, y) values by simultaneously incrementing x and y by small steps proportional to the first derivative of x and y . In our case, the x increment is 1, and the y increment is $dy/dx = m$. Since real variables have limited precision, summing an inexact m repetitively introduces cumulative error buildup and eventually a drift away from a true $\text{Round}(y_i)$; for most (short) lines, this will not present a problem.

Program 3.1

The incremental line scan-conversion algorithm.

```
void Line(int x0, int y0, int x1, int y1, int value)
{
    /* Assumes -1 ≤ m ≤ 1, x0 < x1 */
    int x;           /* x runs from x0 to x1 in unit increments. */
    float dy, dx, y, m;

    dy = y1 - y0;
    dx = x1 - x0;
```

¹ If such a low-level function is not available, the SRG_P_pointCoord function may be used, as described in the SRG reference manual.

```

m = dy / dx;
y = y0;
for (x = x0; x <= x1; x++) {
    WritePixel(x, (int) floor(y + 0.5), value); /*Set pixel to value */
    y += m;                                     /*Step y by slope m */
}
}

```

3.2.2 Midpoint Line Algorithm

The drawbacks of function Line are that rounding y to an integer takes time and that the variables y and m must be real or fractional binary because the slope is a fraction. Bresenham developed a classic algorithm [BRES65] that is attractive because it uses only integer arithmetic, thus avoiding the Round function, and allows the calculation for (x_{i+1}, y_{i+1}) to be performed incrementally—that is, by using the calculation already done at (x_i, y_i) . A floating-point version of this algorithm can be applied to lines with arbitrary real-valued endpoint coordinates. Furthermore, Bresenham's incremental technique may be applied to the integer computation of circles as well, although it does not generalize easily to arbitrary conics. We therefore use a slightly different formulation, the *midpoint technique*, first published by Piteway [PITT67] and adapted by Van Aken [VANA84] and other researchers. For lines and integer circles, the midpoint formulation, as Van Aken shows [VANA85], reduces to the Bresenham formulation and therefore generates the same pixels. Bresenham showed that his line and integer circle algorithms provide the best-fit approximations to true lines and circles by minimizing the error (distance) to the true primitive [BRES77]. Kappel discusses the effects of various error criteria in [KAPP85].

We assume that the line's slope is between 0 and 1. Other slopes can be handled by suitable reflections about the principal axes. We call the lower-left endpoint (x_0, y_0) , and the upper-right endpoint (x_1, y_1) .

Consider the line in Fig. 3.6, where the previously selected pixel appears as a black circle and the two pixels from which to choose at the next stage are shown as unfilled circles. Assume that we have just selected the pixel P at (x_P, y_P) and now must choose between the pixel one increment to the right (called the east pixel, E) or the pixel one increment to the right and one increment up (called the northeast pixel, NE). Let Q be the intersection point of the line being scan-converted with the grid line $x = x_P + 1$. In Bresenham's formulation, the difference between the vertical distances from E and NE to Q is computed, and the sign of the difference is used to select the pixel whose distance from Q is smaller as the best approximation to the line. In the midpoint formulation, we observe on which side of the line the midpoint M lies. It is easy to see that, if the midpoint lies above the line, pixel E is closer to the line; if the midpoint lies below the line, pixel NE is closer to the line. The line may pass between E and NE , or both pixels may lie on one side, but in any case, the midpoint test chooses the closest pixel. Also, the error—that is, the vertical distance between the chosen pixel and the actual line—is always $\leq 1/2$.

The algorithm chooses NE as the next pixel for the line shown in Fig. 3.6. Now all we need is a way to calculate on which side of the line the midpoint lies.

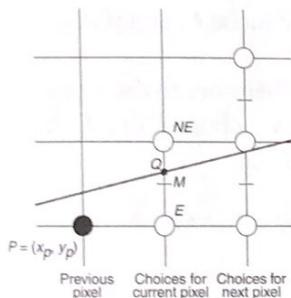


Figure 3.6

The pixel grid for the midpoint line algorithm, showing the midpoint M , and the E and NE pixels to choose between.

Let us represent the line by an implicit function² with coefficients a , b , and c : $F(x, y) = ax + by + c = 0$. (The b coefficient of y is unrelated to the y intercept B in the slope-intercept form.) If $dy = y_1 - y_0$, and $dx = x_1 - x_0$, the slope-intercept form can be written as

$$y = \frac{dy}{dx} x + B;$$

therefore,

$$F(x, y) = dy \cdot x - dx \cdot y + B \cdot dx = 0.$$

Here $a = dy$, $b = -dx$, and $c = B \cdot dx$ in the implicit form.³

It can easily be verified that $F(x, y)$ is zero on the line, positive for points below the line, and negative for points above the line. To apply the midpoint criterion, we need only to compute $F(M) = F(x_P + 1, y_P + \frac{1}{2})$ and to test its sign. Because our decision is based on the value of the function at $(x_P + 1, y_P + \frac{1}{2})$, we define a decision variable $d = F(x_P + 1, y_P + \frac{1}{2})$. By definition, $d = a(x_P + 1) + b(y_P + \frac{1}{2}) + c$. If $d > 0$, we choose pixel NE ; if $d < 0$, we choose E ; and if $d = 0$, we can choose either, so we pick E .

Next, we ask what happens to the location of M and therefore to the value of d for the next grid line; both depend, of course, on whether we chose E or NE . If E is chosen, M is incremented by one step in the x direction. Then,

$$d_{\text{new}} = F(x_P + 2, y_P + \frac{1}{2}) = a(x_P + 2) + b(y_P + \frac{1}{2}) + c,$$

but

$$d_{\text{old}} = a(x_P + 1) + b(y_P + \frac{1}{2}) + c.$$

Subtracting d_{old} from d_{new} to get the incremental difference, we write $d_{\text{new}} = d_{\text{old}} + a$.

We call the increment to add after E is chosen Δ_E ; $\Delta_E = a = dy$. In other words, we can derive the value of the decision variable at the next step incrementally from the value at the current step without having to compute $F(M)$ directly, by merely adding Δ_E .

If NE is chosen, M is incremented by one step each in both the x and y directions. Then,

$$d_{\text{new}} = F(x_P + 2, y_P + \frac{3}{2}) = a(x_P + 2) + b(y_P + \frac{3}{2}) + c.$$

Subtracting d_{old} from d_{new} to get the incremental difference, we write

$$d_{\text{new}} = d_{\text{old}} + a + b.$$

We call the increment to add to d after NE is chosen Δ_{NE} ; $\Delta_{NE} = a + b = dy - dx$.

Let us summarize the incremental midpoint technique. At each step, the algorithm chooses between two pixels based on the sign of the decision variable calcu-

²This functional form extends nicely to the implicit formulation of circles.

³It is important for the proper functioning of the midpoint algorithm to choose a to be positive; we meet this criterion if dy is positive, since $y_1 > y_0$.

lated in the previous iteration; then it updates the decision variable by adding either Δ_E or Δ_{NE} to the old value, depending on the choice of pixel.

Since the first pixel is simply the first endpoint (x_0, y_0) , we can directly calculate the initial value of d for choosing between E and NE . The first midpoint is at $(x_0 + 1, y_0 + \frac{1}{2})$, and

$$\begin{aligned} F(x_0 + 1, y_0 + \frac{1}{2}) &= a(x_0 + 1) + b(y_0 + \frac{1}{2}) + c \\ &= ax_0 + by_0 + c + a + b/2 \\ &= F(x_0, y_0) + a + b/2. \end{aligned}$$

But (x_0, y_0) is a point on the line and $F(x_0, y_0)$ is therefore 0; hence, d_{start} is just $a + b/2 = dy - dx/2$. Using d_{start} , we choose the second pixel, and so on. To eliminate the fraction in d_{start} , we redefine our original F by multiplying it by 2; $F(x, y) = 2(ax + by + c)$. This multiplies each constant and the decision variable (and the increments Δ_E and Δ_{NE}) by 2 but does not affect the sign of the decision variable, which is all that matters for the midpoint test.

The arithmetic needed to evaluate d_{new} for any step is simple integer addition. No time-consuming multiplication is involved. Further, the inner loop is quite simple, as seen in the midpoint algorithm of Prog. 3.2. The first statement in the loop, the test of d , determines the choice of pixel, but we actually increment x and y to that pixel location after updating the decision variable (for compatibility with the circle algorithms). Note that this version of the algorithm works for only those lines with slope between 0 and 1; generalizing the algorithm is left as Exercise 3.2. In [SPRO82], Sproull gives an elegant derivation of Bresenham's formulation of this algorithm as a series of program transformations from the original brute-force algorithm. No equivalent of that derivation for circles has yet appeared, but the midpoint technique does generalize, as we shall see.

Program 3.2

```
void MidpointLine(int x0, int y0, int x1, int y1, int value)
{
    int dx, dy, incrE, incrNE, d, x, y;
    dx = x1 - x0;
    dy = y1 - y0;
    d = dy * 2 - dx;           /* Initial value of d */
    incrE = dy * 2;            /* Increment used for move to E */
    incrNE = (dy - dx) * 2;    /* Increment used for move to NE */
    x = x0;
    y = y0;
    WritePixel(x, y, value);   /* The start pixel */
    while (x < x1) {
        if (d <= 0)           /* Choose E */
            d += incrE;
            x++;
        } else {                /* Choose NE */
            d += incrNE;
            x++;
            y++;
        }
    }
}
```

The midpoint line scan-conversion algorithm.

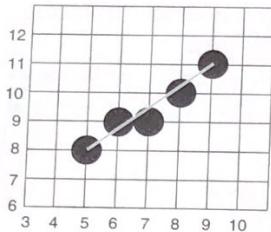


Figure 3.7
The midpoint line from point (5,8) to point (9,11)

```

        }
        WritePixel(x, y, value); /* The selected pixel closest to the line */
    }
}

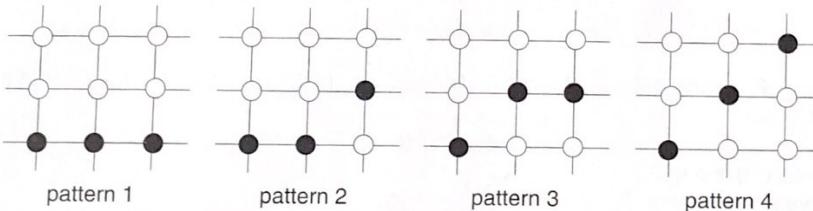
```

For a line from point (5, 8) to point (9, 11), the successive values of d are 2, 0, 6, and 4, resulting in the selection of NE, E, NE, and then NE, respectively, as shown in Fig. 3.7. The line appears abnormally jagged because of the enlarged scale of the drawing and the artificially large interpixel spacing used to make the geometry of the algorithm clear. For the same reason, the drawings in the following sections also make the primitives appear blockier than they look on an actual screen.

Example 3.1

Problem: Develop and program an enhanced version of the midpoint line algorithm that looks ahead two pixels rather than only one.

Answer: An alternative to using the midpoint line algorithm for scan converting lines is the double-step algorithm introduced by Wu and Rokne [WU87]. Like the midpoint algorithm, this technique is incremental; (x_{i+1}, y_{i+1}) can be computed from (x_i, y_i) using only integer arithmetic. The midpoint algorithm first determines the direction (slope) of the line, then it proceeds by using a decision variable to choose between two alternative pixels at each step. The double-step algorithm reduces the number of decisions, i.e., choices for the next pixel, by two. It accomplishes this by looking for the next pair of pixels rather than just the next pixel. Wu [WU87] showed that four patterns can occur:



Wu showed that pattern 1 and pattern 4 cannot occur on the same line. Furthermore, if the slope of the line is greater than $\frac{1}{2}$, pattern 1 cannot occur. Similarly, if the slope of the line is less than $\frac{1}{2}$, pattern 4 can not occur. Thus, by testing the slope, the choice is narrowed to one of three patterns: 1, 2, 3 or 2, 3, 4. Let us look at the case where the slope is between 0 and $\frac{1}{2}$. This excludes pattern 4, as already indicated. A determination needs to be made whether pattern 1, or either pattern 2 or 3 is used to fill in the raster image. The decision variable d is initially set to $d = 4dy - dx$. Then for each increment (step of two raster units), testing whether $d < 0$ determines whether pattern 1 is used. If d is greater than or equal to 0, then a determination must be made between patterns 2 and 3. This is simply the test: $d < 2dy$. To increment d , we use the following rule:

$$d_{i+1} = d_i + 4dy \quad \text{if } d_i < 0 \text{ (pattern 1),}$$

$$d_{i+1} = d_i + 4dy - 2dx \quad \text{otherwise (pattern 2 or 3).}$$

defining primitives could be done a line segment at a time, but that would result in some pixels being drawn that lie outside a primitive's area—see Sections 3.4 and 3.5 for special algorithms to handle this problem. Care must be taken to draw shared vertices of polylines only once, since drawing a vertex twice causes it to change color or to be set to background when writing in **xor** mode to a screen, or to be written at double intensity on a film recorder. In fact, other pixels may be shared by two line segments that lie close together or cross as well. See Exercise 3.7 for a discussion of this issue and of the difference between a polyline and a sequence of connected line segments.

3.3 SCAN CONVERTING CIRCLES

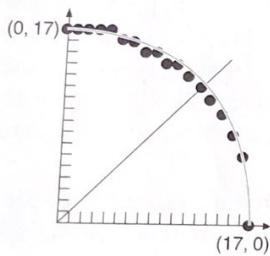


Figure 3.10
A quarter circle generated with unit steps in x , and with y calculated and then rounded. Unique values of y for each x produce gaps.

Although SRGP does not offer a circle primitive, the implementation will benefit from treating the circular ellipse arc as a special case because of its eight-fold symmetry, both for clipping and for scan conversion. The equation of a circle centered at the origin is $x^2 + y^2 = R^2$. Circles not centered at the origin may be translated to the origin by integer amounts and then scan-converted, with pixels written with the appropriate offset. There are several easy but inefficient ways to scan convert a circle. Solving for y in the implicit circle equation, we get the explicit $y = f(x)$ as

$$y = \pm\sqrt{R^2 - x^2}.$$

To draw a quarter circle (the other quarters are drawn by symmetry), we can increment x from 0 to R in unit steps, solving for $+y$ at each step. This approach works, but it is inefficient because of the multiply and square-root operations. Furthermore, the circle will have large gaps (unless R is large) for values of x close to R , because the slope of the circle becomes infinite there (see Fig. 3.10). A similarly inefficient method, that does, however, avoid the large gaps, is to plot $(R \cos\theta, R \sin\theta)$ by stepping θ from 0° to 90° .

3.3.1 Eight-Way Symmetry

We can improve the drawing process of the previous section by taking greater advantage of the symmetry in a circle. Consider first a circle centered at the origin. If the point (x, y) is on the circle, then we can trivially compute seven other points on the circle, as shown in Fig. 3.11. Therefore, we need to compute only one 45° segment to determine the circle completely. For a circle centered at the origin, the eight symmetrical points can be displayed with function CirclePoints (the function is easily generalized to the case of circles with arbitrary origins):

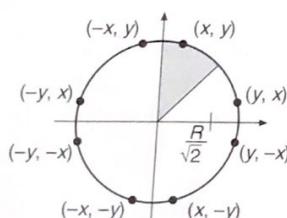


Figure 3.11
Eight symmetrical points on a circle.

```
void CirclePoints (float x, float y, int value);
{
    WritePixel (x, y, value);
    WritePixel (y, x, value);
    WritePixel (y, -x, value);
```

```

    WritePixel (x, -y, value);
    WritePixel (-x, -y, value);
    WritePixel (-y, -x, value);
    WritePixel (-y, x, value);
    WritePixel (-x, y, value);
}

```

We do not want to call CirclePoints when $x = y$, because each of four pixels would be set twice; the code is easily modified to handle that boundary condition.

3.3.2 Midpoint Circle Algorithm

Bresenham [BRES77] developed an incremental circle generator that is more efficient than the methods we have discussed. Conceived for use with pen plotters, the algorithm generates all points on a circle centered at the origin by incrementing all the way around the circle. We derive a similar algorithm, again using the midpoint criterion, which, for the case of integer center point and radius, generates the same, optimal set of pixels. Furthermore, the resulting code is essentially the same as that specified in patent 4,371,933 [BRES83].

We consider only 45° of a circle, the second octant from $x = 0$ to $x = y = R/\sqrt{2}$, and use the CirclePoints function to display points on the entire circle. As with the midpoint line algorithm, the strategy is to select which of two pixels is closer to the circle by evaluating a function at the midpoint between the two pixels. In the second octant, if pixel P at (x_p, y_p) has been previously chosen as closest to the circle, the choice of the next pixel is between pixel E and pixel SE (see Fig. 3.12).

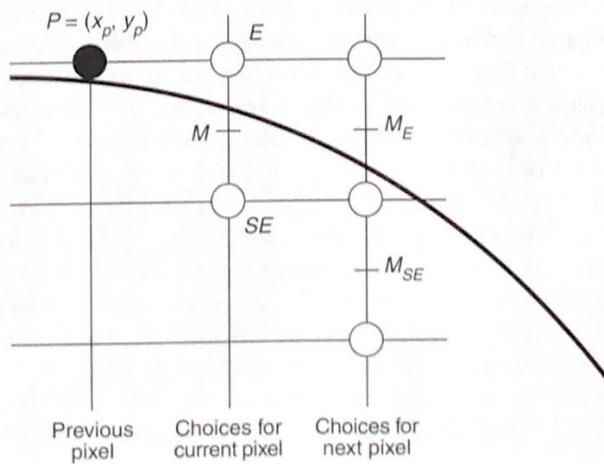


Figure 3.12 The pixel grid for the midpoint circle algorithm showing M and the pixels E and SE to choose between.

Let $F(x, y) = x^2 + y^2 - R^2$; this function is 0 on the circle, positive outside the circle, and negative inside the circle. It can be seen that if the midpoint between the pixels E and SE is outside the circle, then pixel SE is closer to the circle. On the other hand, if the midpoint is inside the circle, pixel E is closer to the circle.

As for lines, we choose on the basis of the decision variable d , which is the value of the function at the midpoint,

$$d_{\text{old}} = F(x_p + 1, y_p - \frac{1}{2}) = (x_p + 1)^2 + (y_p - \frac{1}{2})^2 - R^2.$$

If $d_{\text{old}} < 0$, E is chosen, and the next midpoint will be one increment over in x . Then,

$$d_{\text{new}} = F(x_p + 2, y_p - \frac{1}{2}) = (x_p + 2)^2 + (y_p - \frac{1}{2})^2 - R^2,$$

and $d_{\text{new}} = d_{\text{old}} + (2x_p + 3)$; therefore, the increment $\Delta_E = 2x_p + 3$.

If $d_{\text{old}} \geq 0$, SE is chosen,⁵ and the next midpoint will be one increment over in x and one increment down in y . Then

$$d_{\text{new}} = F(x_p + 2, y_p - \frac{3}{2}) = (x_p + 2)^2 + (y_p - \frac{3}{2})^2 - R^2.$$

Since $d_{\text{new}} = d_{\text{old}} + (2x_p - 2y_p + 5)$, the increment $\Delta_{SE} = 2x_p - 2y_p + 5$.

Recall that, in the linear case, Δ_E and Δ_{SE} were constants; in the quadratic case, however, Δ_E and Δ_{SE} vary at each step and are functions of the particular values of x_p and y_p at the pixel chosen in the previous iteration. Because these functions are expressed in terms of (x_p, y_p) , we call P the point of evaluation. The Δ functions can be evaluated directly at each step by plugging in the values of x and y for the pixel chosen in the previous iteration. This direct evaluation is not expensive computationally, since the functions are only linear.

In summary, we do the same two steps at each iteration of the algorithm as we did for the line: (1) Choose the pixel based on the sign of the variable d computed during the previous iteration, and (2) update the decision variable d with the Δ that corresponds to the choice of pixel. The only difference from the line algorithm is that, in updating d , we evaluate a linear function of the point of evaluation.

All that remains now is to compute the initial condition. By limiting the algorithm to integer radii in the second octant, we know that the starting pixel lies on the circle at $(0, R)$. The next midpoint lies at $(1, R - \frac{1}{2})$, therefore, and $F(1, R - \frac{1}{2}) = 1 + (R^2 - R + \frac{1}{4}) - R^2 = \frac{5}{4} - R$. Now we can implement the algorithm directly, as in Prog. 3.3. Notice how similar in structure this algorithm is to the line algorithm.

The problem with this version is that we are forced to do real arithmetic because of the fractional initialization of d . Although the function can be easily modified to handle circles that are not located on integer centers or do not have integer radii, we would like a more efficient, purely integer version. We thus do a simple program transformation to eliminate fractions.

First, we define a new decision variable, h , by $h = d - \frac{1}{4}$, and we substitute $h + \frac{1}{4}$ for d in the code. Now the initialization is $h = 1 - R$, and the comparison

⁵Choosing SE when $d = 0$ differs from our choice in the line algorithm and is arbitrary. The reader may wish to simulate the algorithm by hand to see that, for $R = 17$, one pixel is changed by this choice.

Program 3.3

The midpoint circle scan-conversion algorithm.

```
void MidpointCircle(int radius, int value)
{
    int x, y;
    float d;

    x = 0;           /*Initialization*/
    y = radius;
    d = 5.0 / 4 - radius;
    CirclePoints(x, y, value);
    while (y > x) {
        if (d < 0) {           /*Select E*/
            d += x * 2.0 + 3;
            x++;
        } else {                /*Select SE*/
            d += (x - y) * 2.0 + 5;
            x++;
            y--;
        }
        CirclePoints(x, y, value);
    }
}
```

$d < 0$ becomes $h < -\frac{1}{4}$. However, since h starts out with an integer value and is incremented by integer values (Δ_E and Δ_{SE}), we can change the comparison to just $h < 0$. We now have an integer algorithm in terms of h ; for consistency with the line algorithm, we will substitute d for h throughout. The final, fully integer algorithm is shown in Prog. 3.4.

Program 3.4

The integer midpoint circle scan-conversion algorithm.

```
void MidpointCircle(int radius, int value)
{
    int x, y, d;

    x = 0;           /*Initialization.*/
    y = radius;
    d = 1 - radius;
    CirclePoints(x, y, value);
    while (y > x) {
        if (d < 0) {           /*Select E*/
            d += x * 2 + 3;
            x++;
        } else {                /*Select SE*/
            d += (x - y) * 2 + 5;
            x++;
            y--;
        }
        CirclePoints(x, y, value);
    }
}
```

Figure 3.13 shows the second octant of a circle of radius 17 generated with the algorithm, and the first octant generated by symmetry (compare the results to Fig. 3.10).

Second-order differences. We can improve the performance of the midpoint circle algorithm by using the incremental computation technique even more extensively. We noted that the Δ functions are linear equations, and we computed them directly. Any polynomial can be computed incrementally, however, as we did with the decision variables for both the line and the circle. In effect, we are calculating **first- and second-order partial differences**, a useful technique that we will encounter again in Chapter 9. The strategy is to evaluate the function directly at two adjacent points, to calculate the difference (which, for polynomials, is always a polynomial of lower degree), and to apply that difference in each iteration.

If we choose E in the current iteration, the point of evaluation moves from (x_P, y_P) to $(x_P + 1, y_P)$. As we saw, the first-order difference is $\Delta_{E\text{old}}$ at $(x_P, y_P) = 2x_P + 3$. Therefore,

$$\Delta_{E\text{new}} \text{ at } (x_P + 1, y_P) = 2(x_P + 1) + 3,$$

and the second-order difference is $\Delta_{E\text{new}} - \Delta_{E\text{old}} = 2$.

Similarly, $\Delta_{SE\text{old}}$ at $(x_P, y_P) = 2x_P - 2y_P + 5$. Therefore,

$$\Delta_{SE\text{new}} \text{ at } (x_P + 1, y_P) = 2(x_P + 1) - 2y_P + 5,$$

and the second-order difference is $\Delta_{SE\text{new}} - \Delta_{SE\text{old}} = 2$.

If we choose SE in the current iteration, the point of evaluation moves from (x_P, y_P) to $(x_P + 1, y_P - 1)$. Therefore,

$$\Delta_{E\text{new}} \text{ at } (x_P + 1, y_P - 1) = 2(x_P + 1) + 3,$$

and the second-order difference is $\Delta_{E\text{new}} - \Delta_{E\text{old}} = 2$. Also,

$$\Delta_{SE\text{new}} \text{ at } (x_P + 1, y_P - 1) = 2(x_P + 1) - 2(y_P - 1) + 5,$$

and the second-order difference is $\Delta_{SE\text{new}} - \Delta_{SE\text{old}} = 4$.

The revised algorithm then consists of the following steps: (1) Choose the pixel based on the sign of the variable d computed during the previous iteration; (2) update the decision variable d with either Δ_E or Δ_{SE} , using the value of the corresponding Δ computed during the previous iteration; (3) update the Δ 's to take into account the move to the new pixel, using the constant differences computed previously; and (4) do the move. Δ_E and Δ_{SE} are initialized using the start pixel $(0, R)$. The revised function using this technique is shown in Prog. 3.5.

Program 3.5

The midpoint circle scan-conversion algorithm using second-order differences.

```
void MidpointCircle(int radius, int value)
{
    /*This function uses second-order partial differences to compute increments*/
    /* in the decision variable. Assumes center of circle is at origin*/
    int x, y, d, deltaE, deltaSE;
```

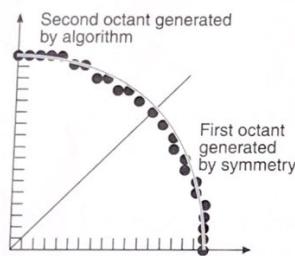


Figure 3.13
Second octant of circle generated with midpoint algorithm, and first octant generated by symmetry.

```

x = 0;           /*Initialization*/
y = radius;
d = 1 - radius;
deltaE = 0;
deltaSE = d - radius * d;
CirclePoints(x, y, value);
while (y > x) {
    if (d < 0) {           /*Select E*/
        d := deltaE;
        deltaE := d;
        deltaSE := d;
    } else {                /*Select SE*/
        d := deltaSE;
        deltaE := d;
        deltaSE := d;
    }
    x++;
    y--;
    CirclePoints(x, y, value);
}

```

8.4 FILLING RECTANGLES

The task of filling primitives can be broken down into two parts: the decision of which pixels to fill (this depends on the shape of the primitive, as modified by clipping), and the easier decision of with what value to fill them. We first discuss filling unclipped primitives with a solid color; we will deal with pattern filling in Section 13.6. In general, determining which pixels to fill consists of taking successive scan lines that intersect the primitive and filling in spans of adjacent pixels that lie inside the primitive from left to right.

To fill a rectangle with a solid color, we set each pixel lying on a scan line running from the left edge to the right edge to the same pixel value; that is, we fill each span from y_{min} to y_{max} . Spans exhibit a primitive's **spatial coherence**: the fact that primitives often do not change from pixel to pixel within a span or from scan line to scan line. We exploit coherence in general by looking for only those pixels of which changes need. For a solidly shaded primitive, all pixels on a span are set to the same value, which provides **span coherence**. The solidly shaded rectangle also exhibits **among-scan-line coherence** in that consecutive scan lines that intersect the rectangle are identical; later, we also use **edge coherence** for the edges of general polygons. We take advantage of various types of coherence not only for scan converting 2D primitives, but also for rendering 3D primitives, as will be discussed in Section 13.4.